

The Samsung Insights Engine. **Delivered.**

The vision is defined and the brief is written. The phases, the use cases and the success criteria already exist, because you built them, and the only cost left is every month spent waiting for the tool itself. This document removes the wait: **the Samsung Insights Engine, live inside Samsung's private cloud, four weeks from signature.**

4 weeks

from signature to a working tool inside Samsung's private cloud

1 quarter

to deliver the brief's 12-month roadmap, Phases 1 to 3

Zero

data leaving Samsung. No third parties, no egress, ever

£10–15k

per month. A licence, not a development project

This is not a description of a platform, it is your tool. The questions below are taken from your brief, so pick one and watch the Insights Engine answer it.

▸ Top 3 washing machine features?

▸ Which touchpoints matter for TV shoppers?

▸ Create a deck: how people shop for a new TV

▸ Did the Currys banner actually work?

▸ Summarize competitor customer profiles

⚠ Try asking about product IP

This is the print edition. The interactive version of this document, where each of these questions can be watched being answered live, is at hey.milo.ai/samsung-insights-engine

https:// insights.samsung.com

00:31.0

S A M S U N G Insights Engine

Insight team workspace · SSO

L

What are the top 3 washing machine features that drive the purchase decision?

✓ Searching the Insights Bank · 1,000 files

✓ Reading 14 relevant decks, including chart images

✓ Cross-checking sources and recency

✓ Writing a cited answer

IE

 Insights Engine

CONFIDENCE · HIGH

Capacity, running cost and quietness lead the decision, in that order.

1. Drum capacity (9kg+). First-ranked driver among family buyers in three of four recent studies. [↩ Laundry Category Drivers UK 2025 · p.14](#)

2. Energy efficiency. Moved from fourth to second since 2023 as running costs entered the decision. [↩ Home Appliance Purchase Journey · p.22](#)

3. Quiet operation. Decisive in open-plan homes, strongest among 25 to 39 year olds. [↩ Living Trends Deep Dive 2024 · p.31](#)

GROUNDING

 Three sources agree on the top three. One older study (2022) ranks price second; it is flagged as pre-dating the energy price rises and weighted accordingly.

That took **31 seconds**. The manual version is a request in the Insight team queue.

Illustrative demonstration built on typical consumer electronics category data, not Samsung's. Once deployed, every answer is generated exclusively from the Insights Bank and the data sources Samsung connects, with citations back to the source document and page. Where no source exists, the Engine says so.

Don't take the demonstration's word for it.

The demonstration above runs on illustrative data, and we would rather you judged the Engine on your own. Hand us three real decks from the Insights Bank under NDA, including the ones where the charts are flat images, and within a week you will watch your own questions answered on your own content, cited and clicked back to source. It costs Samsung nothing but the NDA, and it turns this document's boldest claims into something you have already seen working.

One page, written to be forwarded.

This summary assumes no context, and it can go to anyone in the building, at any level, exactly as it stands.

The opportunity. Samsung's Commercial and Marketing teams hold more data and insight than almost any competitor: a thousand research decks, sell-in and sell-out by retailer, brand trackers, media spend, CRM and device usage data. Very little of it is joined up, so getting an answer out of it takes days of manual work by the Insight team, many requests are repeated, and most marketing briefs go to agencies without insight in them. The Insights Engine brief, written inside Samsung, defines the solution: one trusted, source-grounded tool that any marketer can ask in plain English, with every answer cited back to an approved Samsung source.

The proposal. We deliver that tool. The Samsung Insights Engine runs entirely inside Samsung's private cloud, carries Samsung's name, and is live for users four weeks after signature. It is powered by Milo, a proven insight and analytics engine already deployed in white-label form inside global businesses including Kellanova. There is no development project to fund and no build risk to carry: the brief's 12-month roadmap is delivered in one quarter, and the analytical capability the brief places in Phase 3 is native to the product on day one.

The impact, in three currencies. **Money saved:** fewer repeated research purchases, less agency rework, and no internal build programme to resource. **Time saved:** insight requests answered in seconds rather than days, and the Insight team freed from servicing requests to work on opportunities and threats. **Growth:** joined-up data the business has never been able to analyse together, surfacing patterns across sell-out, brand and media spend that change where the next pound is invested.

The commercial structure. A licence of £10,000 to £15,000 per month depending on tier over a two-year initial term, with a one-off £10,000 deployment fee, hosted by Samsung on Samsung's infrastructure. The first three months operate as the pilot and are judged against the brief's own MVP success criteria. Samsung's data never leaves Samsung, no third-party services are involved, and Samsung retains the deployed application, with software assurance and source-code escrow offered in contract. There is no lower-risk route to this capability, and no faster one.

↓ [Download this summary as a one-page PDF](#)

What Monday morning looks like.

Three people, one tool, the same working week transformed. This is the brief's "example user experience", made concrete.

A MARKETING MANAGER

Writes a brief that lands first time.

As she types a product name and audience into a campaign brief, the Engine reads the keywords and drops in the relevant insight: purchase drivers, audience profile, competitor context, each line cited. The agency gets the rationale built in. **Fewer revision rounds, stronger creative, and she never filed a request.**

AN INSIGHT ANALYST

Stops servicing. Starts investigating.

The requests that used to fill his week now answer themselves, and the answers persist, so they are not asked twice. His Monday goes on the question leadership actually cares about: what is driving the competitor's share gain, and where is the best counter. **The Engine does the number-crunching, and he does the thinking.**

A SENIOR LEADER

Walks into the trading review armed.

Before the meeting she asks for a one-page summary of category performance: share movement, sell-out trend, tracker shifts, each figure cited. What she used to commission a week ahead she now has in the time it takes to read it. **Decisions move at the speed of the question, not the queue.**

The brief's 12-month roadmap. One quarter.

Your brief invites an alternative phasing, so here is ours. We keep your four phases exactly as you defined them and compress the calendar, because the Engine already exists: this is a deployment, not a build.

Weeks 1–4

Phase 1

Document intelligence

Deployed white-label in Samsung's private cloud. SSO connected, existing permissions inherited. The Insights Bank ingested, with automatic ingestion of new files as they are uploaded to SMC. Cited conversational Q&A, curated deck and Word summary outputs, click-back to source, confidence indicators, user feedback facility. **And Phase 3's analytical core is live on day one, because it is the native product, not a later ambition.**

Weeks 5–8

Phase 2

Structured data

Qlik Sense data connected: shares, market values, brand scores, social metrics, sell-in and sell-out. ACE Briefing Portal, Teams and SharePoint integrated via API. Keyword-triggered insight flowing into marketing activity briefs.

Month 3

Phase 3

Analytical intelligence

The full analytical suite tuned and validated on Samsung's own data and decks: shares, growth, averages across time windows, significance testing, correlation, regression, and written trend commentary. **The pilot is judged here, against the brief's own MVP success criteria.**

Month 4+

Phase 4

Predictive and future

Predictive modelling, proactive alerts and auto-generated KPIs rolling out. Voice querying and the audio and avatar presentation layer delivered within your Phase 4 window.

The compression is not effort or optimism: **the engine underneath this tool already exists**, is already deployed inside global businesses, and is already doing this work.

04 · CAPABILITY COVERAGE

Everything the brief asks for. Dated.

Every requirement in the brief, answered against the compressed roadmap. Nothing here is aspiration: each line carries the window in which Samsung users will be using it.

FROM YOUR BRIEF	HOW IT IS DELIVERED	WHEN
Search, synthesise and summarise the Insights Bank (up to 1,000 mixed files), answering only from given sources, saying "I don't know" where no source exists	Core retrieval capability, grounded and abstaining by design	WEEKS 1-4
Automatic ingestion of new files as they are uploaded to SMC	Auto-refresh on file update, a pattern already proven in our Kellanova deployment	WEEKS 1-4
Citations with document titles, page numbers, click-back to source, confidence indicators	Every answer, every format, no exceptions	WEEKS 1-4
Charts that exist only as flat images inside decks	Vision-model extraction reads the data out of chart images; validated on Samsung's real decks during the pilot	WEEKS 1-4
User feedback facility on answer quality	Built in, feeding the tuning loop during the pilot	WEEKS 1-4
Short conversational answers to plain-English questions	The core product	DAY 1
Curated Word summaries and PowerPoint deck outputs	Generated with citations on every slide	WEEKS 1-4
Dashboards, replacing or augmenting Qlik views	Native. The Engine builds dashboards and small applications on your data, not just answers	DAY 1
Qlik Sense data integration	Shares, values, brand scores, social metrics, sell-in and sell-out connected	WEEKS 5-8
ACE Briefing Portal, Teams and SharePoint integrations	API-based, using the platforms' standard interfaces	WEEKS 5-8

FROM YOUR BRIEF	HOW IT IS DELIVERED	WHEN
Keyword-triggered insight inserted into marketing activity briefs	Brief keywords trigger retrieval, synthesis and insertion, with optional supporting slides	WEEKS 5–8
Working with other AI agents	Agent interoperability over MCP, the emerging open standard	YOUR PHASE 2 WINDOW
Phase 3 analysis: market share, growth, averages across time windows, significance testing, correlation, regression, written trend commentary	Native, day one, zero additional development. This is the Engine's bread and butter, and the reason it exists	DAY 1
Auto-generated KPIs triggered by keywords in the marketing brief	Layered onto the brief-support pipeline once it is live	MONTH 3
Predictive modelling: shifts in attitudes, market dynamics, competitive behaviour	Dedicated time-series and statistical models, not LLM guesswork; validated against a paid agency report at 95% agreement (see Evidence)	MONTHS 3–4
Proactive alerts within automated reports	Threshold and anomaly alerts on the measures you nominate	MONTH 4+
Voice querying	Committed roadmap	YOUR PHASE 4 WINDOW
Podcast and audio outputs, avatar-presented decks	Committed roadmap	YOUR PHASE 4 WINDOW
Embedding into a broader architecture	The Engine also runs headless, powering other front-ends over API (detail in the appendix)	DAY 1

Samsung-selected AI models, including Samsung's own

Model-agnostic deployment; internally hosted models, including Gauss, are supported

WEEKS 1-4

THE DISTINCTION THAT MATTERS

Anyone can let you talk to your files. That is a commodity now, and it is the whole of what most vendors will offer you. Almost no one can make an engine genuinely **analyse** your data: correlations, significance, regression, prediction, grounded and cited and explained in plain English. Your brief's Phases 3 and 4 are where the value compounds, and they are where we **start**, not where we hope to arrive.

05 • IMPACT

Measured in Samsung's three currencies.

The brief estimates a 20% efficiency uplift and a 50% effectiveness uplift in selected use cases, using industry benchmarks. We adopt those assumptions rather than inventing better ones, and we commit to replacing them with measured results during the pilot.

MONEY SAVED

Waste stops compounding.

Repeated requests disappear because answers persist, and duplicated research stops being bought twice. Briefs that carry insight from day one mean **fewer agency revision rounds**, and every round avoided is production budget recovered. There is also no build programme to fund, because the engineering is already done.

TIME SAVED

Days become seconds.

A request that takes days of deck-mining is answered, cited, in under a minute. The brief's own 20% efficiency assumption frees **hundreds of Insight team hours a year** for higher-impact work, and marketing decisions stop queuing behind the answers they depend on.

GROWTH

The currency vendors never bring.

Joined-up data is a competitive weapon: when sell-out, brand metrics and media spend are analysed together, patterns surface that no dashboard can show, and they change **where the next pound goes**. That is growth over maintenance, and it is the part only an analytical engine can deliver.

"Did that banner on Currys actually work?"

Every retail team asks it. Almost none can answer it, because the media spend sits in one system, the sell-out sits in another, and no analyst has the week it takes to join them for a single banner. So retailer media is renewed on relationship and instinct, year after year.

The Engine joins them. It compares sell-out in banner weeks against matched non-banner periods, controls for seasonality and promotions, tests whether the difference is statistically significant, and states its confidence and its limitations in plain English. You saw it do this in the demonstration above.

One answer like that, once a quarter, redirects more media spend than this licence costs in a year. That is what "anything that delivers growth more than the maintenance" looks like in practice.

MEASURED IN THE PILOT

Your KPIs, adopted as our commitments.

- ✓ Average time to answer common insight requests
- ✓ Number of repeated insight requests reduced
- ✓ Number of users using the tool monthly
- ✓ Percentage of marketing briefs containing sourced insight
- ✓ Reduction in agency rework or creative revision rounds
- ✓ Insight Team hours freed per month
- ✓ Stakeholder satisfaction score
- ✓ Answer quality rating
- ✓ Percentage of answers with valid citations
- ✓ Number of reused historical insight assets
- ✓ Reduction in duplicated research or duplicated analysis

This list is taken from the brief without amendment. We baseline each measure in week two, report monthly, and the pilot is judged on the movement. No vendor-invented metrics, no goalpost-moving.

Your walls. Your data. Your tool.

The plain-English version. The precise version, for the technical reviewers, is in the appendix. Both say the same thing.

The Samsung Insights Engine is deployed entirely inside Samsung's private cloud, under Samsung's name, and it answers exclusively from sources Samsung provides. It never scrapes the web, never calls a third-party service, and never sends a byte of Samsung data outside Samsung's environment: not to us, not to anyone. Users sign in with their existing Samsung credentials, and the Engine inherits the confidentiality levels, data ownership rules and role restrictions Samsung already enforces.

We hold SOC 2 and ISO 27001 aligned practices, and the appendix covers them. But the honest punchline is this: **certifications matter far less when the tool never leaves your walls.** The strongest security control in this proposal is the architecture itself.

INSIDE THE ENGINE'S SCOPE

Connected, permissioned, cited

- The Insights Bank: research decks, category reports, trackers
- Qlik Sense dashboard data: shares, values, brand and social metrics, sell-in and sell-out
- Marketing and media data: campaign spend, paid, social, PR, influencer
- Approved internal sources: financials, CRM segments, NPS, web and app analytics
- External sources Samsung licenses: syndicated tracking, forecasting, benchmarking

Fenced by design, not by policy

- ✕ **Technical and product IP.** Never connected, never indexed, never answerable.
- ✕ **Personnel and people data.** Same fence, same architecture.
- ✕ The public web. The Engine does not browse or scrape.

07 · EVIDENCE

Proven where it counts.

Two deployments carry the weight of this proposal's two boldest claims: white-label deployment inside a global enterprise, and analysis that stands up against paid professional work.

WHITE-LABEL · GLOBAL FMCG

Kellanova (Kellogg's): the deployment pattern Samsung is asking for, already running.

Milo runs as a **fully white-labelled application** inside Kellanova's own cloud environment, on Databricks and Azure, and their data never leaves it. The knowledge base **refreshes automatically as files update**, which is exactly the SMC auto-ingest pattern in your brief, and role-based access controls who sees what. It was delivered to a demanding, technology-sceptical user group, so what Samsung is asking for is not our first attempt at it.

A European national lottery operator put our analysis against their agency's. It matched.

Their incumbent analytics agency produced a paid predictive report. Milo was given the same underlying data and reproduced the findings with **roughly 95% agreement, in minutes**. That is the standard behind every Phase 3 and Phase 4 claim in this document: the analysis is done by dedicated statistical and time-series models, and it holds up against professional human work, at a fraction of the time and cost.

Breadth: more than ten deployments across regulated gaming, hospitality, retail and FMCG, each running against live commercial data in production. We will gladly arrange reference calls with existing customers as part of your evaluation.

A licence, not a development project.

Licensed over a two-year initial term, billed monthly, with a one-off deployment fee of £10,000 covering the four-week installation. Samsung hosts the application and retains the deployment: the structure itself is the answer to being beholden to a vendor. Two tiers.

£10,000 / month

The brief's Phases 1 to 3, delivered in one quarter, for the Insight team and the HEA and Corporate Marketing teams.

- ✓ **Self-hosted, white-label licence** inside Samsung's private cloud
- ✓ Full document intelligence: Insights Bank, SMC auto-ingest, cited Q&A, deck and summary outputs
- ✓ **Full analytical suite:** shares, growth, significance, correlation, regression, trend commentary
- ✓ Dashboards on your data, built and maintained in the product
- ✓ Qlik Sense, ACE Briefing Portal, Teams and SharePoint integrations
- ✓ Keyword-triggered insight into marketing briefs
- ✓ Standard support and updates included

£15,000 / month

Everything in Core, plus the brief's Phase 4 as it releases, and room for the Engine to grow across the business.

- ✓ **Everything in Core**
- ✓ **Phase 4 capabilities as released:** predictive modelling, proactive alerts, auto-generated KPIs
- ✓ Voice querying, podcast and audio outputs, avatar-presented decks per the Phase 4 window
- ✓ Agent interoperability (MCP) for working with other AI agents
- ✓ **Expanded user population** across additional divisions and teams
- ✓ **Priority support:** a named contact and faster response commitments
- ✓ Priority influence on the product roadmap

Cost by phase, answered directly: there is no per-phase build fee, because there is no build. A one-off deployment fee of £10,000 covers the four-week installation into

Samsung's private cloud, Phases 1 to 3 are covered by the Core licence, Phase 4 by Complete, and the only cost beyond that is Samsung's own cloud infrastructure spend, which is modest and estimated in the appendix.

Ongoing support model: standard support and updates are included at Core, and Complete adds priority support with a named contact and faster response commitments.

Software assurance and escrow: every major Milo capability released during the term is made available to the Samsung deployment, guaranteed in writing. A source-code escrow agreement is also offered in contract, so Samsung can run the Engine independently in any future where we cannot.

09 · THE PILOT

Months one to three, judged by your own test.

The pilot is the first quarter of the contract, at the Core rate.

SUCCESS-GATED · NOT A FREE POC

We do not propose a free proof of concept: free pilots are slow, under-resourced and unaccountable, and you have already lost months. This pilot is the real deployment, fully staffed from day one, with a defined test and a clean exit.

"The Phase 1 MVP should be judged successful if it demonstrates faster access to trusted insight, high-quality cited answers, clear user adoption, and measurable reduction in repeated manual insight requests."

THE SUCCESS TEST, TAKEN VERBATIM FROM YOUR BRIEF

We baseline your KPI list in week two, report against it monthly, and at the end of month three the Engine is judged against that test. **If it fails its own test, Samsung walks away** and keeps everything learned. If it passes, you are not starting a project: you are already live, already adopted, and already into Phase 4.

A short list, deliberately.

The brief asks for resource requirements across the BI, Marketing Ops and IT teams. This is the complete list. No Samsung engineers build anything.

A named contact in BI, Marketing Ops and IT

A few hours per week during month one, tapering after. They review, approve and point; they do not build.

SMC access credentials

Read access to the Insights Bank so ingestion and auto-refresh can begin in week one.

Private cloud environment provisioning

Days, not months. We supply the full specification on day one and assist with the provisioning itself.

SSO application registration

Standard identity-provider setup so users sign in with existing Samsung credentials and permissions.

A security review slot in week one

We bring complete documentation and answer everything in the room. The architecture makes this a short conversation.

Qlik Sense read access, by week five

Needed for Phase 2 only. Nothing in weeks one to four depends on it.

That is the entire ask: **the four-week clock runs from signature**, and the day-level plan for weeks one and two is in the appendix.

Named honestly, each with its answer.

RISK OR DEPENDENCY	MITIGATION
Environment provisioning slips and delays the four-week timeline	The infrastructure specification is delivered on day one, the Engine runs on any Samsung-approved cloud, and we assist with provisioning directly. This is the single dependency that can move the go-live date, which is why it is first on the list and first in the plan.
Chart-image extraction quality varies across Samsung's deck styles	Vision extraction is validated on Samsung's real decks in weeks one to four. Low-confidence extractions are flagged rather than presented as fact, and the user feedback facility routes problem decks into the tuning loop.
Adoption stalls after novelty fades	Adoption is a pilot KPI, taken from your brief, and reported monthly. Onboarding is run per team, the brief-support integration puts the Engine inside existing workflow rather than beside it, and usage data tells us where to intervene early.
Permission mapping is more complex than expected	The Engine inherits Samsung's existing SSO roles and confidentiality levels rather than inventing a parallel scheme. The two fenced data categories are excluded architecturally, which removes the hardest cases from scope entirely.
Qlik Sense or platform API access arrives late	Phase 2 integrations are sequenced for weeks five to eight precisely so that weeks one to four have no external dependencies. A named IT contact and read access by week five keeps Phase 2 on time.

Vendor continuity: what if the supplier disappears?

Samsung retains the deployed application, running in Samsung's own cloud. If we vanished tomorrow, the Engine would still be running where it is, with Samsung's data still inside Samsung's walls. The licence buys updates, support and the roadmap, not our permission for the tool to keep working, and a source-code escrow agreement is offered in contract so that independence is guaranteed, not assumed. We answer this unprompted because it is the right question to ask of any young company.

12 · THE DECISION

You said you would like this yesterday.

If Samsung ultimately decides this tool must be built internally, then this document is still useful to you: it is the blueprint of what good looks like, and we would rather have written it than not. But the brief was written because the need is now, and an internal Phase 1 alone is the better part of a year before the first question gets answered. This proposal is the only route that gets you close to yesterday. The Engine exists, the deployment pattern is proven, and the four weeks can start whenever you say so.



Andreas

Chief Executive. Accountable for every commitment in this document, and on the calls, not above them.

**Philipp**

Delivery. Owns the four-week plan and the deployment inside Samsung's cloud.

**Tom**

Commercial lead. Your first call for anything, from contracting to scope.

Not an account team that hands you off after signature. These are the three people on the calls, from the first provisioning session to go-live and beyond.

**The Milo team**

Prepared for Samsung Electronics UK · The Insights Engine brief

TECHNICAL APPENDIX

For the technical reviewers.

The body of this document is written for its business readers. This appendix is written for the people who will interrogate it: precise, complete, and free of marketing. Section references A1 to A10 are used in the coverage checklist at the end.

Everything inside Samsung's perimeter. Nothing outside it.

The application is deployed in its entirety within Samsung's private cloud. AWS, GCP and Azure are all routine targets, and Samsung's own cloud platform is supported. There is no connection to Milo's servers, no telemetry, no third-party services in the serving path, and no data egress of any kind. The deployment is containerised and delivered with infrastructure-as-code, so Samsung's team can inspect, reproduce and audit the entire footprint.

SMC / Insights Bank

Research decks, reports · auto-ingest on upload

Structured sources

Qlik Sense data, warehouse, internal APIs

External licensed data

Syndicated tracking, forecasts, benchmarks



Ingestion & extraction

Document parsing, vision extraction of chart images, structured connectors

Index & store

Vector and structured indices, permission-tagged at ingestion

Analysis engines

Dedicated statistical and time-series models, not the LLM



Grounded generation

Samsung-selected LLMs, incl. internally hosted (Gauss)

Samsung Insights Engine

Q&A, decks, summaries, dashboards, brief insertion · SSO & RBAC

Integration surface

ACE Briefing Portal, Teams, SharePoint, headless API

NO ARROW CROSSES THIS BOUNDARY. THAT IS THE SECURITY MODEL.

The Engine also operates headless: the same retrieval, analysis and generation services are exposed over an internal API, so the Engine can power other front-ends or sit inside a broader Samsung architecture if one emerges later.

Model-agnostic by design. The LLM is not the analyst.

The deployment supports Samsung-selected language models, including internally hosted models such as Gauss, and commercial models deployed within Samsung's own cloud tenancy. Model choice is a configuration decision Samsung controls, and it can change over the life of the deployment without rework.

The critical design decision is what the LLM is **not** asked to do. Language models are poor arithmetic and statistics engines, and we do not use them as one. All quantitative work, including shares, growth rates, windowed averages, significance tests, correlation, regression and time-series forecasting, is executed by dedicated statistical and modelling components. The LLM's role is language: understanding the question, orchestrating retrieval and analysis, and writing the grounded commentary on results it did not itself compute. This is the direct answer to the reasonable objection that LLMs cannot do real data analysis. They cannot, and in this architecture they are never asked to.

Inherited, not reinvented.

- **Identity:** SSO against Samsung's identity provider. No separate credentials, no shadow user store.
- **Authorisation:** role-based access control mapped to Samsung's existing confidentiality levels, data ownership and user role restrictions, as the brief requires. Row-level security applies to structured data, so two users asking the same question see only what each is entitled to see.
- **Permission-tagged retrieval:** access rules are attached to content at ingestion and enforced at query time, before generation. An answer cannot cite a source its reader is not permitted to open.
- **Fenced categories:** technical and product IP, and personnel and people data, are excluded architecturally. No connector to those systems exists in the deployment, no such content is indexed, and no configuration change inside the application can widen the fence.
- **Audit:** queries, sources consulted and outputs are logged within Samsung's environment for review.
- **Practice:** our engineering and operations follow SOC 2 and ISO 27001 aligned controls, and we operate for clients in regulated industries, including licensed gaming. Under full self-hosting, however, the certifications are largely moot: Samsung's existing cloud controls govern the running system, and no Milo personnel can reach Samsung data.

Provenance by design, not by promise.

Unstructured content follows an ingestion, enrichment, vectorisation and retrieval pipeline; structured data is queried directly. At answer time the Engine retrieves candidate sources under the user's permissions, and generation is constrained to retrieved material. Every answer carries its citations: document title, page number where relevant, and click-back to the source.

Confidence indicators are computed from source agreement, source recency and retrieval strength, and are shown with the answer rather than hidden in a log.

Where sources are missing, the Engine says "I don't know" instead of improvising. Where sources are incomplete, outdated or contradictory, it says that too, identifies the gap, and states what additional evidence or human review is required, exactly as the brief specifies. Hallucination is addressed structurally rather than rhetorically: unsupported generation is excluded by the grounding constraint, quantitative claims come from the analysis engines rather than the language model, and the citation requirement makes every claim checkable by its reader. The user feedback facility closes the loop, routing weak answers into tuning.

This appendix intentionally describes the pipeline at architecture level rather than component level. We are happy to go deeper under NDA in a technical session.

Built for exactly the estate the brief describes.

- **Unstructured:** PowerPoint, PDF, Word and Excel, at the Insights Bank's scale of up to 1,000 mixed files and well beyond. Slide structure, tables and embedded objects are parsed; charts that exist only as flat images are read by vision-model extraction, which recovers the underlying series and labels so the data inside them becomes citable and analysable. Extraction quality is validated on Samsung's real decks during weeks one to four, and low-confidence extractions are flagged, not asserted.
- **Automatic ingestion:** the Engine watches the SMC location and ingests new and updated files as they arrive, with no manual step. This auto-refresh pattern is already running in our Kellanova deployment.
- **Structured:** warehouse tables, Qlik Sense data, and internal APIs connect through standard connectors. The Qlik scope in the brief (market shares, market values, brand awareness, consideration and purchase intent, social metrics, sell-in and sell-out) is read directly, and Phase 3 adds automated ingestion of further internal and external datasets.

Named, because the brief names them.

Samsung Marketing Cloud (SMC) for the Insights Bank and auto-ingest. Qlik Sense for structured marketing data. ACE Briefing Portal for keyword-triggered insight insertion into activity briefs. Microsoft Teams and SharePoint for delivery of answers and summaries where users already work. Beyond the named list, the headless API covers anything else in the estate that can speak HTTP, which in practice is everything.

The startup-mortality answer, in writing.

Samsung retains the deployed application in Samsung's own environment. If Milo ceased to exist, the Engine would continue to run exactly where it runs today, on Samsung's infrastructure, with Samsung's data. The licence covers what a living vendor provides: updates, new capability as the roadmap releases, support (standard at Core, priority with a named contact at Complete), and provisioning and tuning assistance.

Two further protections are offered in contract. **Software assurance:** every major Milo capability released during the term is made available to the Samsung deployment, so the branch Samsung runs never falls behind the product. **Source-code escrow:** the code is held by an independent escrow agent and released to Samsung should Milo cease trading or be acquired, so Samsung can operate the Engine independently in any future where we cannot.

Four weeks, with the first two at day level.

- **Days 1–2:** kick-off; infrastructure specification handed over; environment provisioning begins with our assistance; SSO application registration requested; security review scheduled.
- **Days 3–5:** application deployed into the provisioned environment; SSO connected; smoke tests pass; security review held with full documentation.
- **Days 6–8:** SMC connected with read credentials; initial ingestion of the Insights Bank; vision extraction pass over chart-image decks; permission mapping applied and verified.
- **Days 9–10:** first cited answers validated side-by-side with the Insight team; feedback facility switched on; tuning loop begins.
- **Weeks 3–4:** user onboarding by team; answer quality tuning against real questions; go-live to the pilot user population; KPI baseline captured in week two, reported from week four.
- **Weeks 5–12:** Phase 2 integrations (Qlik, ACE Briefing Portal, Teams, SharePoint), brief-support pipeline, analytical validation, pilot judgement at the end of month three.
- **Samsung resource:** as listed in the body: named contacts in BI, Marketing Ops and IT for a few hours per week in month one; SMC and Qlik read credentials; provisioning and SSO actions above. No Samsung development effort at any point.

Samsung's own cloud spend, and modest.

Because the Engine runs in Samsung's cloud accounts, infrastructure is billed by Samsung's provider at Samsung's negotiated rates, not by us. At Phase 1 and 2 scale (the pilot user population and an Insights Bank of up to 1,000 files) the expected footprint is a small number of application and database instances plus model serving, typically in the region of £1,500 to £3,000 per month depending on the cloud, the model choice and whether an internally hosted model is used. We provide the precise sizing sheet with the day-one infrastructure specification, and we tune the footprint down as usage patterns settle.

All seventeen, mechanically checkable.

The brief lists seventeen items a response must include. The table below maps each to where it is answered, so evaluation can be a tick exercise.

#	REQUIRED ITEM, FROM THE BRIEF	WHERE IT IS ANSWERED
1	Recommended technical architecture	✓ Appendix A1, A2
2	Recommended MVP scope in Phase 1, and phasing and roadmap if different	✓ Section 03 (the compressed roadmap), Section 09 (pilot scope), A8
3	Data ingestion approach	✓ Appendix A5
4	Approach to unstructured and structured data	✓ Appendix A5; capability table, Section 04
5	Security and permissions model	✓ Section 06 (plain English), Appendix A3 (precise)
6	AI model / RAG approach	✓ Appendix A2, A4
7	Approach to citations and confidence scoring	✓ Appendix A4; demonstrated live in the interactive example
8	Approach to hallucination reduction	✓ Appendix A4
9	Example user experience	✓ The interactive demonstration; Section 02 (Monday morning)
10	Any integration requirements	✓ Appendix A6; Section 10
11	Implementation timeline	✓ Section 03; Appendix A8 (day level, weeks one and two)
12	Cost estimate by phase	✓ Section 08 (phases mapped to tiers), Appendix A9 (infrastructure)
13	Resource requirements from Samsung (BI, Marketing Ops, IT)	✓ Section 10; Appendix A8
14	Risks and dependencies	✓ Section 11
15	Pilot proposal	✓ Section 09
16	Examples of relevant previous work	✓ Section 07

#	REQUIRED ITEM, FROM THE BRIEF	WHERE IT IS ANSWERED
17	Ongoing support model required	✓ Section 08 (included in licence), Appendix A7



Powered by Milo

White-labelled for Samsung Electronics UK